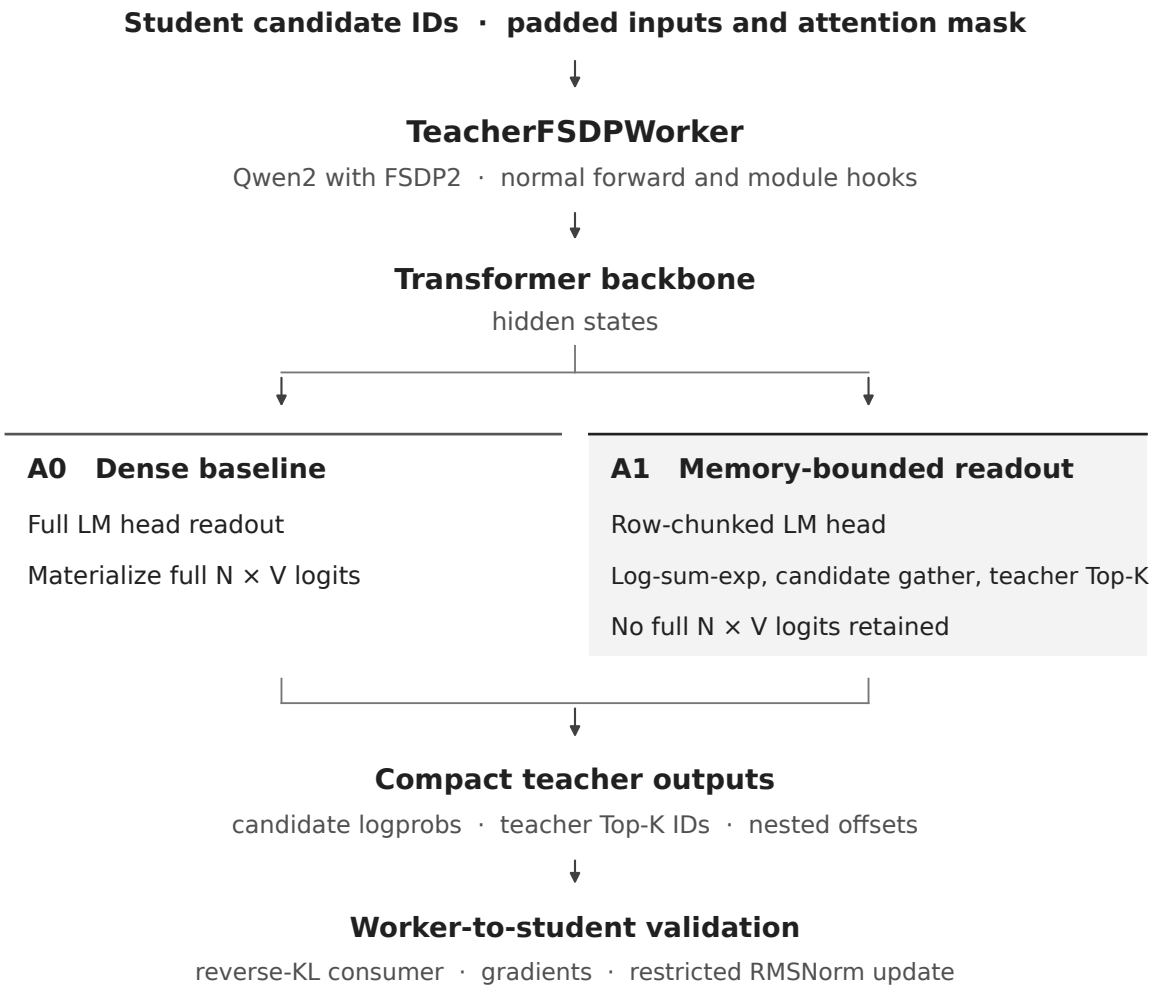
Student candidate IDs · padded inputs and attention mask
TeacherFSDPWorker
Qwen2 with FSDP2 · normal forward and module hooks
Transformer backbone
hidden states
A0 Dense baseline
Full LM head readout
Materialize full N × V logits
A1 Memory-bounded readout
Row-chunked LM head
Log-sum-exp, candidate gather, teacher Top-K
No full N × V logits retained
Compact teacher outputs
candidate logprobs · teacher Top-K IDs · nested offsets
Worker-to-student validation
reverse-KL consumer · gradients · restricted RMSNorm update