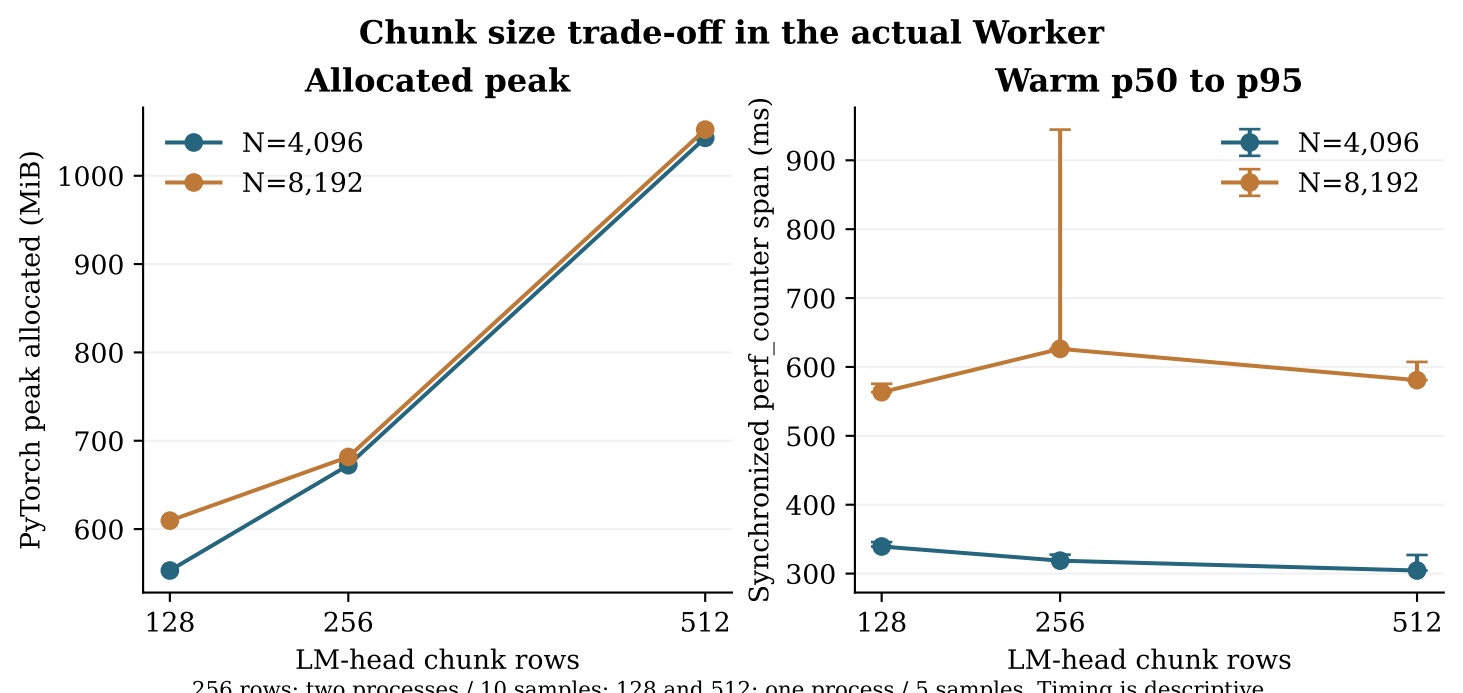
Chunk size trade-off in the actual Worker
Allocated peak
PyTorch peak allocated (MiB)
N=4,096
N=8,192
1000
900
800
700
600
128
256
512
LM-head chunk rows
Warm p50 to p95
Synchronized perf_counter span (ms)
N=4,096
N=8,192
900
800
700
600
500
400
300
128
256
512
LM-head chunk rows
256 rows: two processes / 10 samples; 128 and 512: one process / 5 samples. Timing is descriptive.