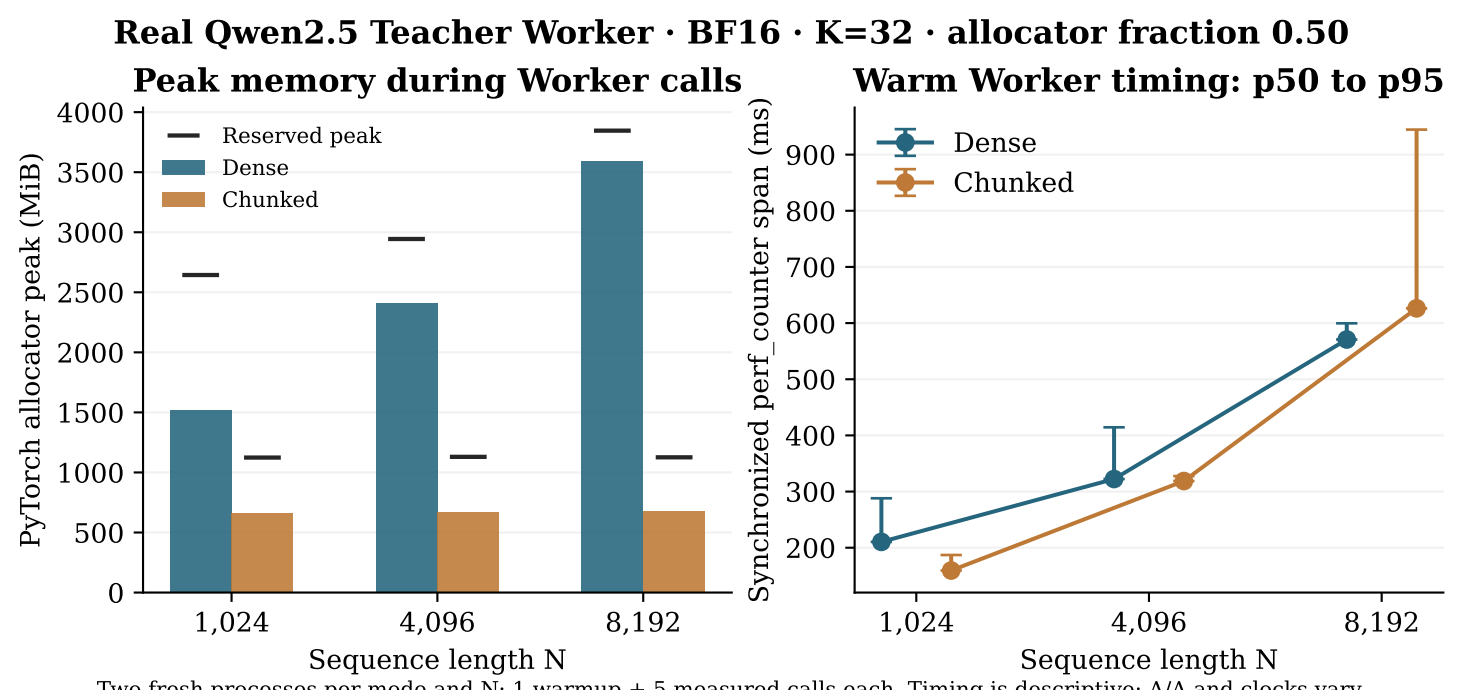
Real Qwen2.5 Teacher Worker · BF16 · K=32 · allocator fraction 0.50
Peak memory during Worker calls
PyTorch allocator peak (MiB)
Reserved peak
Dense
Chunked
0
500
1000
1500
2000
2500
3000
3500
4000
1,024
4,096
8,192
Sequence length N
Warm Worker timing: p50 to p95
Synchronized perf_counter span (ms)
Dense
Chunked
200
300
400
500
600
700
800
900
1,024
4,096
8,192
Sequence length N
Two fresh processes per mode and N; 1 warmup + 5 measured calls each. Timing is descriptive: A/A and clocks vary.